



Classification using Decision Tree

CSE-454: Data Warehousing and
Data Mining Sessional

MD. JAKARIA

LECTURER

DEPT. OF CSE, MIST

Classification

Definition: In machine learning, classification refers to a predictive modeling problem where a class label is predicted for a given example of input data.

Examples of classification problems include:

- Given an example, classify if it is spam or not.
- Given a handwritten character, classify it as one of the known characters.
- Given recent user behavior, classify as churn or not.

There are perhaps four main types of classification tasks that you may encounter; they are:

- Binary Classification
- Multi-Class Classification
- Multi-Label Classification
- Imbalanced Classification

Decision Tree Classification

Definition: A *decision tree* represents a function that takes as input a vector of attribute values and returns a “decision” — a single output value.

As an example, we will build a decision tree to decide whether to wait for a table at a restaurant. The list of attributes that we will consider:

- | | |
|---------------------|--------------------------|
| 1. <i>Alternate</i> | 6. <i>Price</i> |
| 2. <i>Bar</i> | 7. <i>Raining</i> |
| 3. <i>Fri/Sat</i> | 8. <i>Reservation</i> |
| 4. <i>Hungry</i> | 9. <i>Type</i> |
| 5. <i>Patrons</i> | 10. <i>Wait Estimate</i> |

Note that every variable has a small set of possible values.

Decision Tree Classification

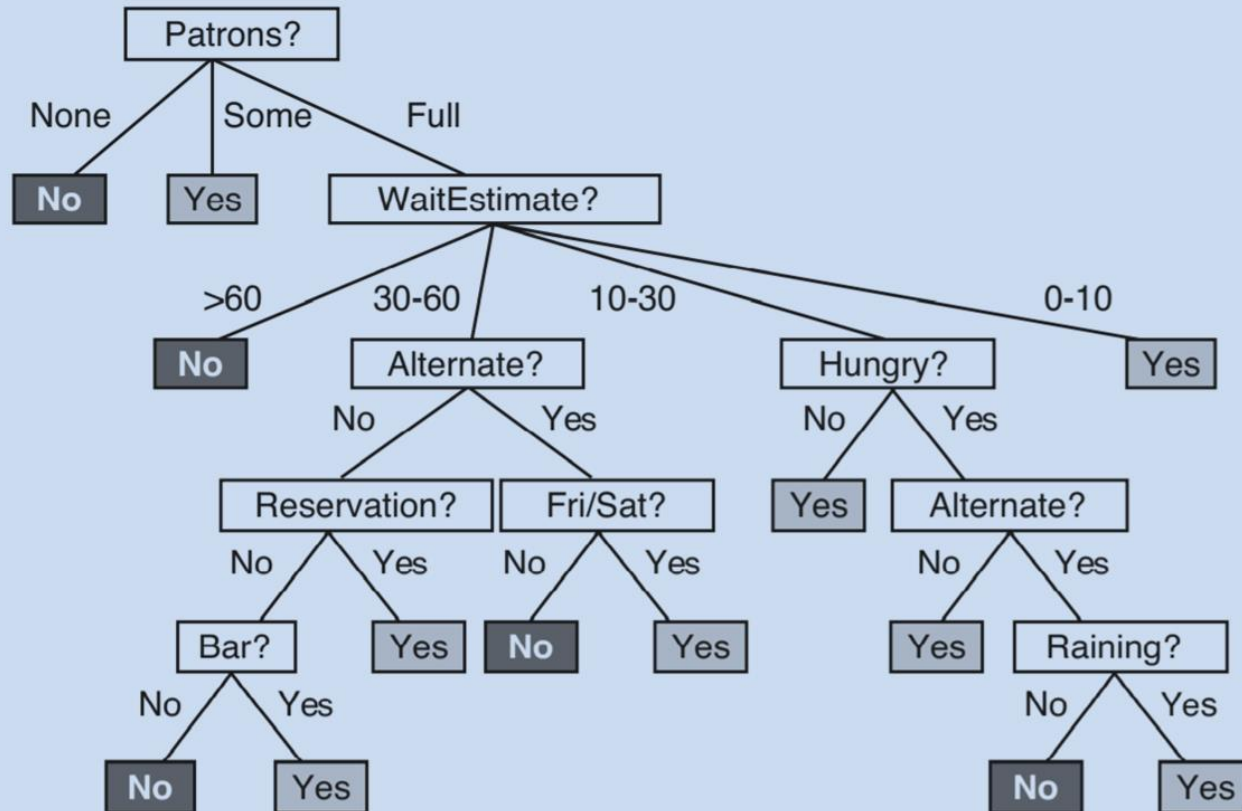


Fig: A decision tree for deciding whether to wait for a table.

Decision Tree Classification

Decision tree -

- A flow-chart-like tree structure
- Internal node denotes a test on an attribute
- Branch represents an outcome of the test
- Leaf nodes represent class labels or class distribution

Use of decision tree:

- Classifying an unknown sample
- Test the attribute values of the sample against the decision tree

Expressiveness of decision trees

A Boolean decision tree is logically equivalent to the assertion that the goal attribute is true if and only if the input attributes satisfy one of the paths leading to a leaf with value true.

$$Goal \Leftrightarrow (Path_1 \vee Path_2 \vee \dots)$$

where each *Path* is a conjunction of attribute-value tests required to follow that path.

$$Path = (Patrons = Full \wedge WaitEstimate = 0-10)$$

Constructing decision trees

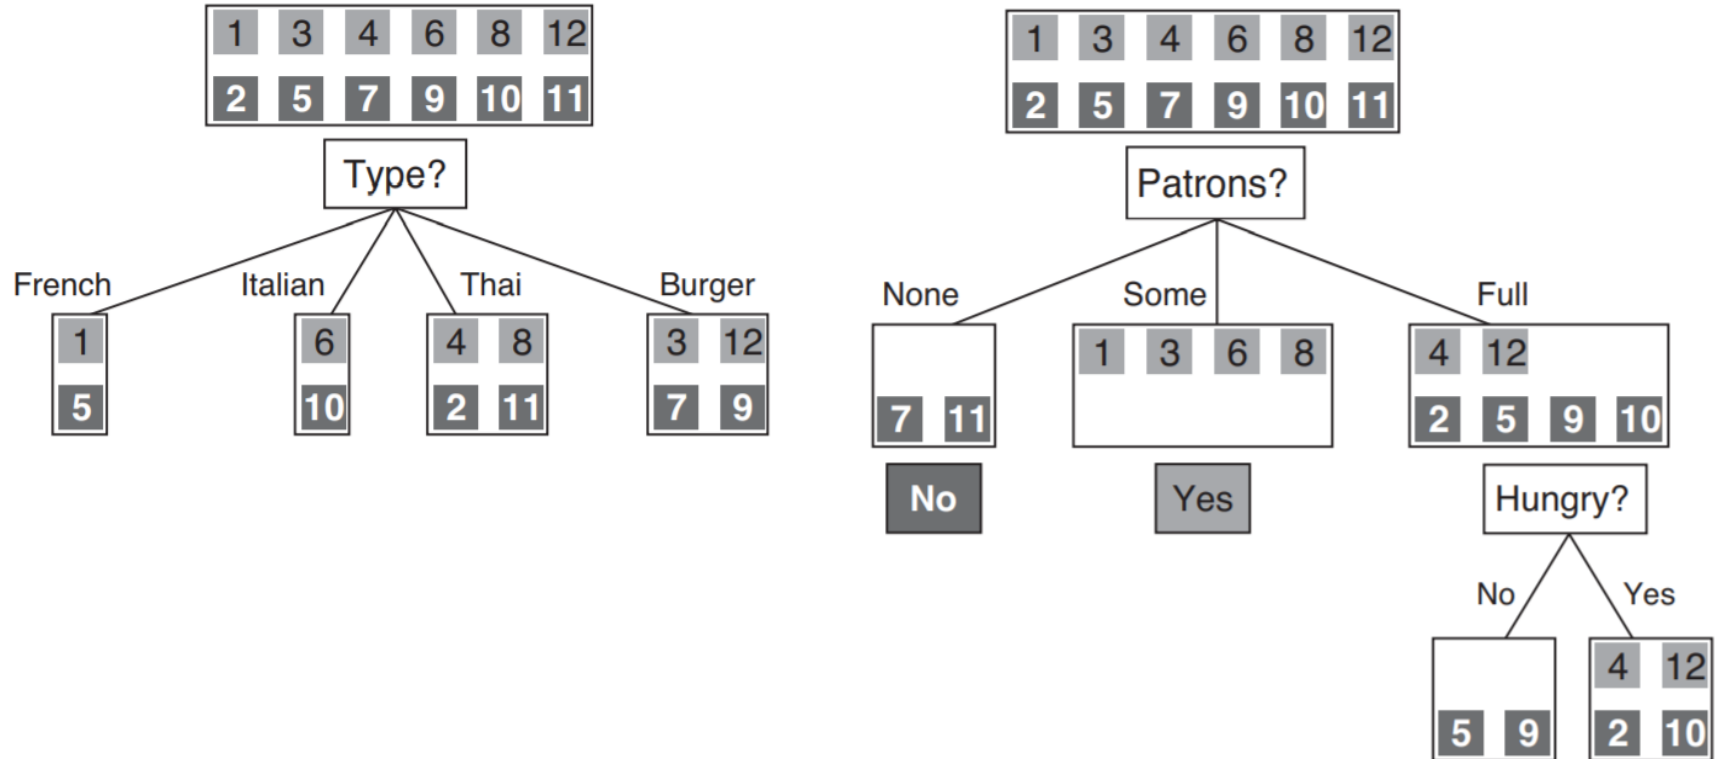
Example	Input Attributes										Goal
	<i>Alt</i>	<i>Bar</i>	<i>Fri</i>	<i>Hun</i>	<i>Pat</i>	<i>Price</i>	<i>Rain</i>	<i>Res</i>	<i>Type</i>	<i>Est</i>	<i>WillWait</i>
x₁	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>French</i>	<i>0–10</i>	<i>y₁ = Yes</i>
x₂	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Thai</i>	<i>30–60</i>	<i>y₂ = No</i>
x₃	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Some</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Burger</i>	<i>0–10</i>	<i>y₃ = Yes</i>
x₄	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Thai</i>	<i>10–30</i>	<i>y₄ = Yes</i>
x₅	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>Full</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>French</i>	<i>>60</i>	<i>y₅ = No</i>
x₆	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$</i>	<i>Yes</i>	<i>Yes</i>	<i>Italian</i>	<i>0–10</i>	<i>y₆ = Yes</i>
x₇	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>None</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Burger</i>	<i>0–10</i>	<i>y₇ = No</i>
x₈	<i>No</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$</i>	<i>Yes</i>	<i>Yes</i>	<i>Thai</i>	<i>0–10</i>	<i>y₈ = Yes</i>
x₉	<i>No</i>	<i>Yes</i>	<i>Yes</i>	<i>No</i>	<i>Full</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Burger</i>	<i>>60</i>	<i>y₉ = No</i>
x₁₀	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>Italian</i>	<i>10–30</i>	<i>y₁₀ = No</i>
x₁₁	<i>No</i>	<i>No</i>	<i>No</i>	<i>No</i>	<i>None</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Thai</i>	<i>0–10</i>	<i>y₁₁ = No</i>
x₁₂	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Burger</i>	<i>30–60</i>	<i>y₁₂ = Yes</i>

Table: Examples for the restaurant domain.

Constructing decision trees

- We want a tree that is consistent with the examples and is as small as possible
- It is an intractable problem to find the smallest consistent tree; there is no way to efficiently search through the 2^{2^n} trees
- With some simple heuristics, we can find a good approximate solution: a small (but not smallest) consistent tree
- The DECISION-TREE-LEARNING algorithm adopts a greedy divide-and-conquer strategy: always test the most important attribute first.
- By “most important attribute,” we mean the one that makes the most difference to the classification of an example.
 - *Type* is a poor attribute
 - *Patrons* is a fairly important attribute

Constructing decision trees



Constructing decision trees

Basic steps (a greedy approach)

- Tree is constructed in a top-down recursive divide-and-conquer manner
- At start, all the training examples are at the root
- Attributes are categorical (if continuous-valued, they can be discretized in advance)
- Examples are partitioned recursively based on selected attributes.
- Test attributes are selected on the basis of a heuristic or statistical measure (e.g., information gain)

Constructing decision trees

Conditions for stopping partitioning

- All samples for a given node belong to the same class
- There are no remaining attributes for further partitioning – majority voting is employed for classifying the leaf
- There are no samples left

The DECISION-TREE-LEARNING algorithm

There are four cases to consider

- i. If the remaining examples are all positive (or all negative), we can answer *Yes* or *No*.
- ii. If there are some positive and some negative examples, then choose the best attribute to split them.
- iii. If there are no examples left, we return a default value calculated from the plurality classification of parent node.
- iv. If there are no attributes left, but both positive and negative examples, return the plurality classification of the remaining examples.

The DECISION-TREE-LEARNING algorithm

function DECISION-TREE-LEARNING (*examples*, *attributes*, *parent_examples*) returns
a tree

if *examples* is empty **then return** PLURALITY-VALUE(*parent_examples*)
else if all examples have the same classification **then return** the classification
else if *attributes* is empty **then return** PLURALITY-VALUE(*examples*)
else

$A \leftarrow \operatorname{argmax}_{a \in \text{attributes}} \text{IMPORTANCE}(a, \text{examples})$

tree \leftarrow a new decision tree with root test *A*

for each value v_k of *A* **do**

$\text{exs} \leftarrow \{ e : e \in \text{examples} \text{ and } e.A = v_k \}$

subtree \leftarrow DECISION-TREE-LEARNING(*exs*, *attributes* – *A*, *examples*)

add a branch to *tree* with label (*A* = v_k) and subtree *subtree*


return *tree*

Choosing attribute tests


- The goal is to have the resulting decision tree as small as possible (Occam's Razor)
 - But, finding the minimal decision tree consistent with the data is NP-hard
- The recursive algorithm is a greedy heuristic search for a simple tree, but cannot guarantee optimality.
- The main decision in the algorithm is the selection of the next attribute to condition on.
- Which is the best attribute?
 - The one which will result in the smallest tree
 - Heuristic: choose the attribute that produces the “purest” nodes
- Need a good measure of purity!
 - Maximal when?
 - Minimal when?

Impurity/Entropy (informal)


Measures the level of impurity in a group of examples



Very Impure Group



Less Impure Group



Minimum Impure Group


Impurity/Entropy (informal)

Measures the level of impurity in a group of examples

$$Entropy = - \sum_i p_i \log_2 p_i$$

- Here, p_i is the probability of class i .
- Compute it as the proportion of class i in the set.

$$Entropy = -\left(\frac{16}{30} \log_2 \frac{16}{30} + \frac{14}{30} \log_2 \frac{14}{30}\right)$$
$$= 0.996$$




Impurity/Entropy (informal)

Measures the level of impurity in a group of examples

$$Entropy = - \sum_i p_i \log_2 p_i$$

- Here, p_i is the probability of class i .
- Compute it as the proportion of class i in the set.

$$Entropy = -\left(\frac{12}{14} \log_2 \frac{12}{14} + \frac{2}{14} \log_2 \frac{2}{14}\right)$$
$$= 0.591$$




Impurity/Entropy (informal)

Measures the level of impurity in a group of examples

$$Entropy = - \sum_i p_i \log_2 p_i$$

- Here, p_i is the probability of class i .
- Compute it as the proportion of class i in the set.

$$Entropy = -\left(\frac{12}{12} \log_2 \frac{12}{12} + \frac{0}{12} \log_2 \frac{0}{12}\right)$$
$$= 0$$

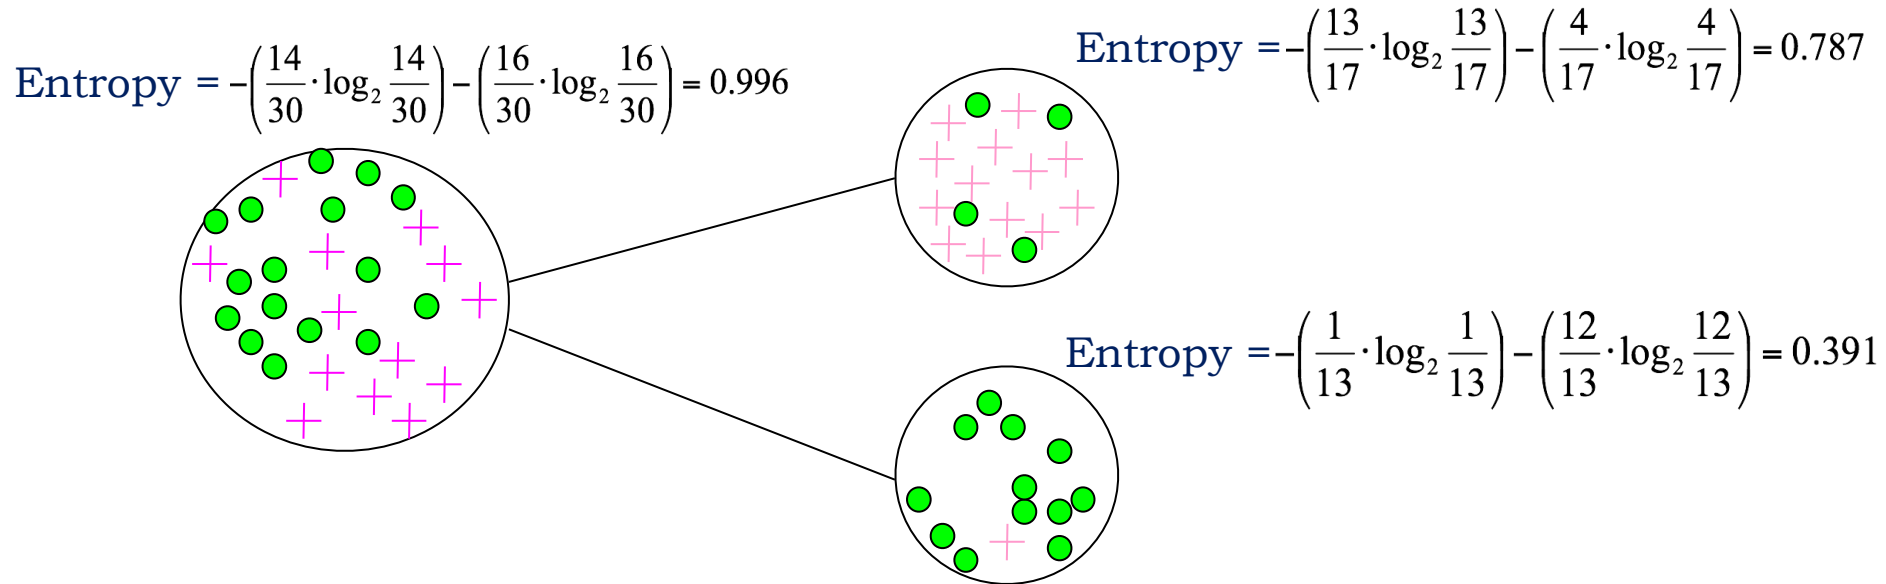


Information Gain

- We want to determine which attribute in a given set of training feature vectors is most useful for discriminating between the classes to be learned.
- Information gain tells us how important a given attribute of the feature vectors is.
- We will use it to decide the ordering of attributes in the nodes of a decision tree.

$$\text{Information Gain} = \text{Entropy}(\text{parent}) - [\text{Average Entropy}(\text{children})]$$

Information Gain




$$(\text{Weighted}) \text{ Average Entropy of Children} = \left(\frac{17}{30} \cdot 0.787\right) + \left(\frac{13}{30} \cdot 0.391\right) = 0.615$$


$$\begin{aligned} \text{Information Gain} &= \text{Entropy}(\text{parent}) - [\text{Average Entropy}(\text{children})] \\ &= 0.996 - 0.615 = 0.38 \end{aligned}$$

Information Gain

Which test is more informative?



Split over whether Balance exceeds 50K



Split over whether applicant is employed

The DECISION-TREE-LEARNING algorithm

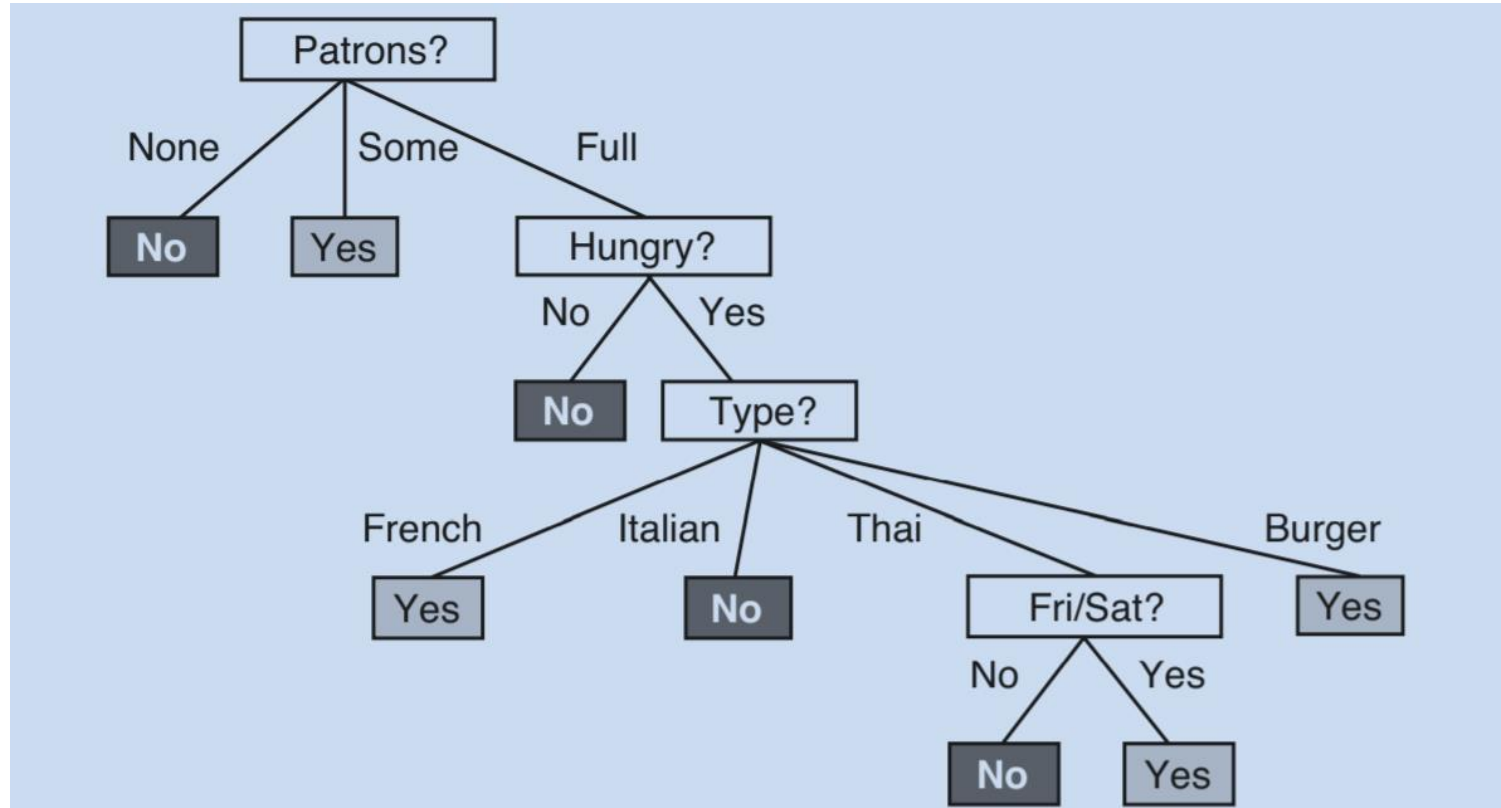


Fig: The decision tree induced from the 12-example training set.

References:

- i. **Artificial Intelligence: A Modern Approach** by Peter Norvig and Stuart J. Russell
Chapter 18 (18.3 Learning Decision Tree)